

---

# Deep Learning and Determining Duplicate Questions

---

**Jacob L. Fenger**  
Oregon State University  
fengerj@oregonstate.edu

**Spike Madden**  
Oregon State University  
email

**Chongxien Chen**  
Oregon State University  
email

## Abstract

This document contains the approach Team Jacob utilized to train a classifier for detecting whether questions are duplicates. We utilized Word2Vec embeddings for vectorizing the questions as well as Keras with a Tensorflow backend for training.

## 1 Introduction

Quora Question Pairs is a Kaggle competition requiring challengers to use advanced techniques to determine if pairs of questions are duplicates. You can find the challenge at:

*<https://www.kaggle.com/c/quora-question-pairs>*

Challengers will need to utilize natural language processing as well as deep learning in order to create a good classification. Our team had no previous experience with natural language processing and during our implementation of the challenge, we learned a lot. There are many aspects to consider when converting sentences to number vectors. This process is otherwise known as vectorization.

Our first approach got a score of while our second approach got a score of. Explanations of these approaches are outlined in the following sections.

## 2 Approach #1

The first approach was influenced by a user named Lystdo who posted a kernel for the challenge. Our approach for the challenge was heavily influenced by his mainly because we had no experience with the natural language processing side of things.

### 2.1 Text Pre-Processing

To ensure accurate results, we pre-processed the text. Initially, this just involved removing all punctuation as well as making everything lowercase. This is necessary because many questions, especially those in the test set, containing words with weird capitalisation or spelling. Our knowledge regarding natural language processing is very limited, but we had confidence that punctuation in the questions provide no additional meaning.

Later on in the implementation of this challenge, we realized that some questions involved contractions, slang terms, or spelling mistakes. We tried our best to handle those by doing replacements with commonly known contractions/slang terms, but some still may have gotten through due to the large data set. The testing set was generated by Quora and many of the questions had spelling mis-

takes. A user called 'Currie32' had an interesting discussion post regarding the cleaning of text for the challenge.

## **2.2 Feature Extraction**

After reading all of the data and pre-processing the text, a necessary step is to analyze the text and generate numeric feature vectors because the majority of machine learning algorithms require them. This process is otherwise known as vectorization.

We first tokenized the questions and then padded the sequences so they were all the same size. Using a pre-trained word2vec model is very helpful in the vectorization of things. There are several pre-trained models in existence. Word2Vec by Google and GloVe by Stanford are just two examples. We decided to use Word2Vec for our implementation. There is a neat package that helps with the interface to Word2Vec called Genism which is what we used as well.

## **2.3 Training**

We utilized Keras and Tensorflow to build a long short term model (LSTM) for the classification. This consisted of an embedding layer involving an embedding matrix generated from the Word2Vec vocab. Other layers included dense and dropout layers with batch normalization in-between. Batch normalization helps with the normalization of data between layers (In short).

## **2.4 Results**

The score we got with this approach was about After updating the text pre-processing with contraction expansion, we got reduced the log-loss by several points.

## **3 Approach #2**

## **4 Comparison of Results**

## **5 Conclusion**

## **Individual Contributions**